

# Quantum-in-the-Loop VLAM Alignment and Certified Control Safety

a quantum curriculum selector in RL alignment, a 22  $\mu$ s certified wrap on the action output, and a disturbance class computed on 128 qubits

Seven guideline sections · every prior-work number carries a receipt

Public repository: <https://github.com/sharadbachani-oss/merlin-quantum-vw-gic2026>

---

## Summary

**−35.4%**

training rollouts to target vs the named GRPO baseline (3 seeds, mean  $\pm$  SD, 71 receipted device jobs) with zero tuned parameters and half the seed variance; +9.3 points safer driving at equal budget, replicated on a second machine

**78→100%**

episodes safe when the uncertified learned driver is wrapped by the certified runtime-assurance monitor — 22  $\mu$ s per decision, envelope computed before measurement and confirmed inside the predicted interval

**0.60 rad/s**

certified boundary under the measured collective disturbance spectrum (128 qubits, two devices) — where white-noise validation over-certifies at every amplitude and the AR surrogate reads 25% safe at the true boundary

## Why the numbers can be trusted

Every claim carries a pre-registration frozen before execution,  $\geq 3$  seeds with mean  $\pm$  SD, in-job controls and cloud job IDs; key results are replicated on an independent machine; the six-arm ladder and rank sweep isolate the quantum component on both tracks; every arm, including the tuned classical Gibbs sampler, is reported in the table. A reviewer replays every headline with `python verify.py`, no credentials.

This is a Phase-1 concept proposal in the seven guideline sections. The Phase-2 PoC puts the same curriculum selector on an accepted VLA backbone (OpenVLA-OFT / LIBERO) and attaches a fidelity certificate and ion-trap replication to the disturbance spectrum.

---

Concept proposal follows (§1–§7, appendices A–B) ·

<https://github.com/sharadbachani-oss/merlin-quantum-vw-gic2026>

# Concept Proposal

## 1. Problem framing

A 7B–10B VLAM does not fit vehicle or robot-controller compute without a training, compression or safety story, and the field's two named blockers are inference latency of 500–2,000 ms against a 100–300 ms control budget and hallucinated actions inconsistent with visual input (DriveVLM-RL, arXiv:2603.18315; VLAM survey, arXiv:2512.16760). Both have one deployment consequence: the VLAM cannot be the thing that decides, so it must be wrapped by something certified and fast, and its action head must be aligned at a sample cost an OEM can afford.

We put a quantum component into the two VLAM stages that decide whether the backbone can ship. **Alignment:** a quantum annealer inside every GRPO iteration selects the rollout curriculum — measured –35.4% rollouts to target versus the named GRPO baseline, three seeds. **Safety:** a certified runtime-assurance monitor (22  $\mu$ s, 4,500 $\times$  inside the budget) wrapping an untrusted policy, whose certified envelope is validated against a **disturbance class computed on quantum hardware** — the real-frequency collective spectrum of an interacting mobility lattice at 128 qubits. Classical validation models misjudge that boundary from both sides (white noise over-certifies at every amplitude; an AR surrogate reads 25% safe at the true boundary). The disturbance spectrum is a real-time many-body quantity, the class of computation every 2025–26 verified quantum-advantage candidate occupies, produced here as a certification input rather than a benchmark.

## 2. Technical approach

**Paradigm.** Hybrid: photonic entropy computing (QCi Dirac-3) for curriculum sampling in the training loop; gate-model real-time dynamics (IBM Heron, 128 qubits) for the disturbance spectrum; classical tensor-network Lyapunov certification and a classical monitor on the vehicle.

**Alignment (scored track).** Each GRPO iteration submits the 64-candidate rollout pool (start  $\times$  disturbance metadata) to Dirac-3, which selects the 16-rollout curriculum under a frontier-plus-diversity objective before any simulation is spent. The objective is encoded on the annealer's native flow manifold (level sets, no penalty walls) — a design rule predicted before running and confirmed on the identical objective: –26% objective, 3.5 $\times$  tighter spread over 12 paired jobs. Quantum work is train-time only.

**Control safety (scored track).** Tensor-network Lyapunov monitor  $V(x) = ||L \cdot \phi(x)||^2$  (3-level tensor-product features, rank 3 of 9), grid-plus-Lipschitz certified, wrapping an uncertified behaviour-cloned policy with covariate shift and 100 ms delay; verified fallback engages at  $V > 0.75c$  with hysteresis. Only monitor and fallback face certification (ISO 26262 / IEC 62061); the learned driver stays a swappable black box.

**Disturbance class on quantum hardware.** 64-rung (128-qubit) interacting lattice on IBM Heron, Trotterised depth series  $k = 0..15$ , 32,768 shots per circuit, 2,656 two-qubit gates per circuit at  $k = 8$ , native heavy-hex with zero SWAPs; equal-depth differential estimator that cancels common-mode decoherence (signal recovered where naive retention is  $10^{-33}$ , cells at  $8.9\sigma$  and  $11.5\sigma$ ,  $0.2\sigma$  preparation control); real-frequency spectrum by direct transform of the real-time series; dominant collective line 0.0625 cycles per step reproduced on two devices.

**Plant.** Kinematic-bicycle lane keeping (lateral/heading error,  $|\delta| \leq 0.25$ ,  $v = 25$  m/s) under crosswind and impulses — the control substrate a VLAM commands. Phase 2 puts the same curriculum selector on an accepted VLA backbone (OpenVLA-OFT or  $\pi$  on LIBERO), see §5.

## 3. Feasibility and resource requirements

| resource                                   | status                                                                                                                         |
|--------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------|
| RL evaluator, six-arm ladder, safety plant | in the public repo; <code>python verify.py</code> (numpy) replays every headline                                               |
| Dirac-3                                    | unmetered allocation; 71 curriculum jobs receipted                                                                             |
| IBM Heron                                  | Startup Program (applied); the spectrum job exists (daa9pn4e74ec73akj9i0); re-flights $\leq 33$ circuits $\times$ 32,768 shots |

| resource                   | status                                                                                               |
|----------------------------|------------------------------------------------------------------------------------------------------|
| VLA backbone run (Phase 2) | OpenVLA-OFT / LIBERO on a single 8-GPU node; 3 seeds; GPU-hours declared in the resource declaration |
| Classical                  | 32-core workstation for the tensor-network certifier and the adversarial classical attack            |

Assumptions: the plant is a control surrogate for the VLA action head; disturbance amplitudes in rad/s rms.  
 Constraints: quantum hardware is never an in-vehicle component; all quantum work is train- or certification-time.

## 4. Expected impact

**RL alignment — rollouts to target (seeds 21/22/23, mean  $\pm$  SD):** A, GRPO-32 random (named baseline)  $586.7 \pm 98.8$ ; **C, Dirac-3-16 untuned  $378.7 \pm 52.9$  (–35.4%)**; D, classical optimiser of the same objective 0/3 (curriculum collapse); E, tuned classical Gibbs sampler  $298.7 \pm 106.9$  with per-task temperature tuning and 2 $\times$  the device's seed variance; C', device on optimiser-grade encoding 0/3. The device's measured advantages are specific: it beats the named baseline by 35.4%, needs zero tuned parameters, and carries roughly half the seed variance; the mechanism is isolated in the ladder (only near-optimal *stochastic* selection trains). At equal budget the curriculum yields a safer driver:  $76.4 \pm 0.7\%$  vs  $67.1 \pm 1.3\%$  safe on the held-out stress suite (**+9.3 points**), replicated on a second machine, with the variance collapse (baseline SD 22–28 points vs curriculum SD 1.8) that de-risks the training run.




Figure 2 — The six-arm alignment ablation (rollouts to target, 3 seeds): random selection fails, a deterministic optimiser of the same objective collapses training, the device and the tuned Gibbs sampler train; the device does it with zero tuned parameters and half the variance.

**Control safety:** policy alone  $77.9 \pm 25.0\%$  safe at the certified boundary  $\rightarrow$   **$100.0 \pm 0.0\%$**  wrapped (3 seeds  $\times$  500 episodes per condition); guard idles under benign conditions (1–12% intervention); 22  $\mu$ s per decision. The safe disturbance margin was computed before measurement (holds to 1.0 rad/s; collapse predicted in [1.2, 1.6]) and the measured collapse landed inside it, on two machines under two perturbation protocols. Versus tuned heuristic clipping: the envelope is known a priori, zero tuned parameters, and the  $V > 0$ ,  $dV \leq 0$  evidence chain an ISO 26262 case consumes. Rank ablation:  $r = 3$  interior-optimal at  $0.462 \pm 0.030$  certified area; the TN certificate covers  $1.57 \pm 0.29\times$  the best quadratic's area.

**Certification under the measured disturbance class (quantum input to the scored track).** Three ensembles at identical rms power, 3 seeds  $\times$  1,000 episodes per cell: against the **measured collective spectrum** the certified boundary is **0.60 rad/s rms** ( $95.6 \pm 0.8\%$  safe;  $98.4 \pm 0.3\%$  at 0.50). White-noise validation grades  $100.0 \pm 0.0\%$  safe at every amplitude through 1.3 rad/s where the true safe rate is 39% — unbounded over-certification; the AR(1) surrogate reads  $25.2 \pm 2.1\%$  at the certified boundary itself. A shallow 16-point anchor-sector spectrum places the boundary at 0.88 rad/s, a 32% envelope error only the full-depth computation catches. The vehicle runs the 22  $\mu$ s monitor; quantum computes the disturbance model at certification time.




Figure 1 — Safe-episode rate of the certified plant against disturbance amplitude under three ensembles at identical rms power: white noise, AR(1) surrogate, and the measured 128-qubit collective spectrum. 3 seeds  $\times$  1,000 episodes per cell.

**Resource allocation.** A pre-registered compression study established that INT4 already meets the challenge bar classically ( $3.97\times$  at 2.45% accuracy drop vs  $\geq 2\times$  at  $\leq 5\%$ ) and that the remaining mixed-precision space is flat — the device solved its allocation objective exactly (DP gap 0.0, 9/9 instances) and beat random 5–25 $\times$ ; the finding tells an OEM where not to spend quantum budget, which is why this entry concentrates on the two stages with measured headroom.

### Quantum advantage — the frontier wall and the crossing, stated and bounded

**The field's direction, and its wall.** Quantum reinforcement learning's advantage results — amplitude-encoded policy evaluation (arXiv:2505.11862), quantum natural policy gradient at  $O(\epsilon^{-1.5})$  (arXiv:2501.16243), Boltzmann-machine sampling (arXiv:2511.04856) — are proofs of concept for fault-tolerant machines; none reports an empirical advantage on existing devices, and a quantum sampler in a training loop has not beaten a tuned classical sampler on hardware. Our six-arm ladder maps that wall exactly: the device and a tuned Gibbs sampler both train, a deterministic optimiser and random selection do not — so the sampler is where the mechanism is proven, not where the quantum contribution is scored.

**The crossing.** We put the quantum computation where it is not a sampling task: the **disturbance class the safety certificate is validated against** is a real-time collective spectrum of a 128-qubit interacting lattice,  $k = 0..15$ , 2,656 two-qubit gates at  $k = 8$ . It is adjudicated from both sides: our own tensor-network attack on the shipped data reproduces the hardware to  $0.4\sigma$  at  $k = 4$  and within  $1.4\sigma$  across five wavevectors at  $k = 6$ , then fails past  $k = 8$  (operators exceeding  $2 \times 10^8$  terms); the hardware cells beyond that boundary stand at  $8.9\sigma$  and  $11.5\sigma$  with a  $0.2\sigma$  preparation control, and the dominant line is reproduced on a second device. In the customer's units the crossing is the 0.60 rad/s certified boundary that white-noise validation over-certifies at every amplitude and the AR surrogate misjudges by 70 points. **Where it is decisive:** past  $k = 8$ , beyond the reach of our own commissioned tensor-network attack (a measured boundary for that method, not a classical impossibility result — a Pauli-propagation attack on the same observable is the next comparator), with  $k \leq 6$  serving as the classically verified anchor of the same instrument; Phase 2 adds the fidelity certificate and the ion-trap replication.

**Scalability to industrial relevance.** The curriculum selector scales with the rollout pool (64 today; 949 variables at the device maximum, equal-structure templates rebind angles per iteration), and the disturbance-class lattice tiles by rung on the same heavy-hex fabric (64 rungs today; 78 on Heron r3, more on Nighthawk's square lattice without changing the compilation). Quantum work is train- and certification-time only, so fleet size does not enter the per-vehicle cost; the 22  $\mu\text{s}$  monitor runs on existing controller silicon.

**Business value, bounded.** In an OEM's own units: (i) alignment cost — 35.4% fewer rollouts to target against the GRPO baseline means, at a declared cost per rollout  $r$  and  $n$  alignment runs per model release, a saving of  $\approx 0.35 \cdot 586 \cdot r \cdot n$ ; (ii) certification — a disturbance envelope that is 32% wrong (shallow spectrum) or 70 points wrong at the operating point (AR surrogate) is either a field incident or an over-conservative envelope that leaves performance on the table; the measured class prices that margin at 0.60 rad/s exactly; (iii) the certified wrap lets the VLAM be upgraded without re-certifying the safety case, which is where the recurring cost sits. Stated as

functions of the OEM's cost inputs, not as fixed sums.

**What a successful PoC demonstrates.** The same curriculum selector inside the alignment loop of an accepted VLA backbone with a  $\geq 10\%$  rollout reduction at equal reward, 3 seeds, paired CI above zero against the strongest classical sampler at equal total cost; and a certified control envelope for that backbone's action output, validated against a hardware-computed disturbance class carrying a fidelity certificate.

5. Validation plan

Pre-registered: reward, robustness metric, seeds, total evaluation budget, disturbance amplitudes, tolerance. **Alignment:** OpenVLA-OFT (or  $\pi$ ) on LIBERO, backbone/demonstrations/reward evaluator/initial checkpoints/training budget fixed; arms = native curriculum, tuned classical Gibbs, original encoding, random, best deterministic rule; samplers crossed with the same optimiser; paired tasks and independent seeds; scores = held-out shifted-task success and instruction consistency, constraint violations reported separately; report error at equal budget and cost to a common target; acceptance = paired CI above zero within the same full training/evaluation cost. **Safety:** 3 seeds  $\times$  500–1,000 episodes per cell, envelope sealed before measurement, heuristic-clipping threshold sweep as baseline. **Disturbance class:** re-read the 128-qubit series under the instrument (atlas layout, even-sector encoding, predicted pair T2, visibility inversion); attach a parity-syndrome fidelity bound; replicate the dominant line on an ion-trap device; report the commissioned tensor-network attack (reproduces hardware to  $0.4\sigma$  at  $k = 4$ , fails past  $k = 6$ ) and a Pauli-path run at matched accuracy. Every result: frozen pre-registration,  $\geq 3$  seeds with mean  $\pm$  SD, in-job controls, cloud job IDs, negatives kept.

6. Hybrid / cross-domain integration

Quantum work is train/validation-time only. Alignment: Dirac-3 returns index sets that are verified classically before use. Safety: certificate candidates pass the deterministic grid-plus-Lipschitz certifier; the vehicle runs the 22  $\mu$ s monitor and fallback. Disturbance class: computed on Heron at certification time, read by inverted estimators, consumed by the certifier as an ensemble. The runtime-assurance pattern makes the uncertified VLAM swappable without re-certification — the ISO 26262 / IEC 62061 posture an OEM can consume. Devices are operated as instruments: idle-ZZ spectrometry against a 24-day atlas (9/9 kingston, 8/9 fez edges), pair T2 predicted from the atlas to 7–9%, native-operator energies recovered to 0.006 on 5/5 live tiles, computation placed in even-ZZ sectors the dominant noise channel cannot see.

7. Team capability

Merlin Quantum is the quantum division of Merlin Digital (50+ technology FTE): Suhail Bachani (Founder & CEO, Principal Investigator), Dr. Hiro Bachani PhD (Program Director), Rohit Bachani (co-founder), Mitul Sawlani (engineering, Purdue), Mohamed Jafrun (engineering), Zeena Furtado (finance & operations), Roshan Bhairwani (financial services & deep tech, London), Dr. Ana Baroni MSc (domain specialist). GIC 2026 dual-track finalist. Programme record: 259 receipted QPU jobs, 12.3 million shots; largest circuits 128 qubits and 2,720 two-qubit gates; structured error law  $C(D) = \exp(-0.00903D - 3.96 \times 10^{-5} D^2)$  measured on-device ( $5.1\times$  vs the white-noise extrapolation); five-way independent error rate 0.47%/gate; in-kind CHSH anchor  $S = 2.3604 (+35.9\sigma, replica +37.0\sigma)$ . Published automotive quantum work is annealing-based routing, materials and scheduling; to our knowledge no published work certifies a control envelope against a hardware-computed disturbance class.

8. Scope, with treatment

(i) The plant is a lane-keeping surrogate, not a VLA backbone — the OpenVLA-OFT/LIBERO run is the PoC's first milestone. (ii) A per-task-tuned classical Gibbs sampler reaches fewer raw rollouts — the device's scored claim is the named baseline at  $-35.4\%$ , zero tuned parameters and half the seed variance. (iii) Compression headroom is flat beyond INT4 — quantum budget is deliberately not spent there. (iv) The disturbance spectrum carries no fidelity certificate yet — the parity-syndrome bound and ion-trap replication are Phase-2 deliverables. (v) The certified envelope is validated on the stated plant and disturbance classes — a new vehicle or fleet condition is re-validated, not assumed.

### Appendix A — Hardware job register (every result regenerates from archived counts)

| measurement                                                                | machine | job id(s)            |
|----------------------------------------------------------------------------|---------|----------------------|
| 128q jam-relaxation spectrum v2, $k = 0..15$ ,<br>33 $\times$ 32,768 shots | ibm_fez | daa9pn4e74ec73akj9i0 |

| measurement                                                                     | machine                | job id(s)                                                                                           |
|---------------------------------------------------------------------------------|------------------------|-----------------------------------------------------------------------------------------------------|
| 156q real-time collective interface, exact-theorem anchors 0.9836 / 0.9806      | ibm_fez                | d9rm4j9dsedc73agrb70;<br>SCout<br>d9rm47opdb6s73e53kqg                                              |
| Two-sided classical boundary series, 16 × 32,768                                | ibm_kingston           | d9rdfb1dsedc73agh5ng                                                                                |
| Idle-ZZ atlas, 9/9 edges                                                        | ibm_kingston           | da9l3t1qtnsc73d1nhd0,<br>da9l9rkjbipc73ff0aqq                                                       |
| Idle-ZZ atlas, 8/9 edges                                                        | ibm_fez                | da9eoe6rbfbs73chiq0g,<br>da9lqeerbfbbs73chq63g                                                      |
| Native-operator E■ inverted, 5/5 tiles                                          | ibm_kingston / ibm_fez | da9vgsmrbfbs73ci44d0,<br>da9vblkjbp73ffaio0                                                         |
| Directed-path traffic card — converges classically, kept as negative control    | ibm_kingston           | results/trafficD_kingston_result.json                                                               |
| CHSH in-kind anchor $S = 2.3604 (+35.9\sigma)$ , same-day replica $+37.0\sigma$ | ibm_fez                | GIC ledger<br>chsh_*_20260805_*.json                                                                |
| Flow-encoding test, 6 paired trials (12 jobs)                                   | QCi Dirac-3            | results/flow_encoding_test.json                                                                     |
| RL curriculum arm C, vw_rl_C_s21.json (20 jobs)                                 | QCi Dirac-3            | 6a76595c08442f441bbb5bd2<br>...<br>6a765b2208442f441bbb5be5<br>— full list in the public repository |
| RL curriculum arm C, vw_rl_C_s22.json (23 jobs)                                 | QCi Dirac-3            | 6a765b8708442f441bbb5be6<br>...<br>6a765d4508442f441bbb5bfc<br>— full list in the public repository |
| RL curriculum arm C, vw_rl_C_s23.json (28 jobs)                                 | QCi Dirac-3            | 6a765d5b08442f441bbb5bfd<br>...<br>6a765f6b08442f441bbb5c18<br>— full list in the public repository |
| Dirac-trained certificate, trackb_seed21.json (9 jobs)                          | QCi Dirac-3            | 6a783b7f08442f441bbb5d37<br>...<br>6a783be408442f441bbb5d3f<br>— full list in the public repository |
| Dirac-trained certificate, trackb_seed22.json (9 jobs)                          | QCi Dirac-3            | 6a783a7c08442f441bbb5d25<br>...<br>6a783ad108442f441bbb5d2d<br>— full list in the public repository |
| Dirac-trained certificate, trackb_seed23.json (9 jobs)                          | QCi Dirac-3            | 6a783adf08442f441bbb5d2e<br>...<br>6a783b3608442f441bbb5d36<br>— full list in the public repository |
| Safety, envelope, rank ablation, cert-gap sweeps                                | CPU (two machines)     | vw_rta_demo.json,<br>vw_lyap_*.json, results/vw_cert_gap_*_v2.json                                  |

### ### Appendix B — Resource declaration and notes

Classical: 32-core workstation, no GPU used for the delivered results; simulation environment numpy/qiskit versions pinned in the repo. Quantum: Dirac-3 (71 + 12 jobs), IBM Heron fez/kingston (shot counts per job in receipts). Phase-2 VLA run: single 8-GPU node, GPU-hours to be declared. Deadline runway: Phase I closes

2026-09-15.

### Appendix C — Claim ledger (measured · planned · comparator · quantum attribution · cost · acceptance)

| claim                                                | status                        | classical comparator                                                              | quantum attribution                  | total cost charged                  | acceptance threshold                             |
|------------------------------------------------------|-------------------------------|-----------------------------------------------------------------------------------|--------------------------------------|-------------------------------------|--------------------------------------------------|
| RL alignment –35.4% rollouts vs GRPO baseline        | measured (3 seeds, 71 jobs)   | tuned Gibbs 298.7 (per-task tuning), random, optimiser                            | Dirac-3 sampler                      | 71 jobs ≈ minutes QPU               | ≥ 10% at equal reward, paired CI > 0             |
| Same selector on OpenVLA-OFT / LIBERO                | planned (Phase-2 milestone 1) | strongest classical curriculum, tuned Gibbs                                       | sampler                              | 8-GPU node, 3 seeds, GPU-h declared | ≥ 10% rollouts-to-target, CI > 0                 |
| Certified boundary 0.60 rad/s under measured class   | measured (3 seeds × 1,000)    | white-noise, AR(1) surrogates                                                     | 128q spectrum as certification input | 1 job × 33 circuits                 | boundary inside predicted interval               |
| 128q spectrum beyond classical attack ( $k \geq 8$ ) | measured; certificate planned | our TN attack ( $0.4\sigma$ at $k=4$ , fails $k \geq 8$ ); Pauli-propagation next | hardware                             | 33 × 32,768 shots                   | parity-syndrome fidelity bound + ion replication |
| 78 → 100% safe under certified wrap                  | measured (classical)          | tuned clipping                                                                    | none                                 | CPU                                 | envelope sealed before measurement               |